

Parallel & Distributed Processing II:

parallel processing on manycore chips

Using CUDA

Eric Aubanel
Winter 2010, UNB Fredericton

Use NVIDIA's resources

- ▶ **In Blackboard**
 - ▶ Getting Started Guide
 - ▶ Programming Guide
 - ▶ Best Practices Guide
 - ▶ Reference Manual
- ▶ **CUDA Zone:** www.nvidia.com/object/cuda_home.html
- ▶ Many other resources on the Web

Compiling and Running

- ▶ Set environment variables:

```
export PATH=/usr/local/cuda/bin:$PATH
export DYLD_LIBRARY_PATH=/usr/local/cuda/lib:
    $DYLD_LIBRARY_PATH
```

- ▶ Put combined CPU/GPU code in one `.cu` file

- ▶ Compile

```
$ nvcc -o prog prog.cu
```

- ▶ Can use gcc flags such as `-O`

- ▶ Optimization of sequential code important for fair comparison

- ▶ Run

- ▶ Same as running any CPU executable

```
$ prog arg1 arg2 ...
```

Measuring Performance

► Elapsed time of CPU code

```
double wallClockTime() { //time in seconds
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return (1000000*tv.tv_sec+tv.tv_usec)/1.0e6;
}

double time = wallClockTime();
// code to be timed
double time = wallClockTime() - time;
```

Measuring Performance

- ▶ Can use CPU timer to time GPU code, but...
 - ▶ Some CUDA API functions are asynchronous
 - ▶ Kernel calls
 - ▶ Memory copies with **Async** suffix
 - ▶ To avoid this problem use `cudaThreadSynchronize()` immediately before starting and stopping timer
 - ▶ Blocks CPU until all CUDA calls previously issued have completed

Measuring Performance

► CUDA GPU timer

- Uses GPU clock, so more accurate timings

```
cudaEvent_t start, stop;  
float time;  
cudaEventCreate(&start);  
cudaEventCreate(&stop);  
  
cudaEventRecord( start, 0 );  
kernel<<<grid,threads>>> ( p1, p2, ... );  
cudaEventRecord( stop, 0 );  
cudaEventSynchronize( stop ); /* block until event  
    actually recorded */  
  
cudaEventElapsedTime( &time, start, stop );  
cudaEventDestroy( start ); cudaEventDestroy( stop );
```

Measuring Device Bandwidth

- ▶ “Bandwidth is one of the most important gating factors for performance. Almost all changes to code should be made in the context of how they affect bandwidth.”
 - *CUDA 2.3 Best Practices Guide*
- ▶ Compare theoretical and effective bandwidth
 - ▶ Theoretical
 - ▶ Can calculate (see best practices guide), but running `/Developer/GPU Computing/C/bin/darwin/release/bandwidthTest` will give actual achievable DRAM bandwidth

Effective bandwidth

- ▶ $[(\text{bytes read per kernel} + \text{bytes written per kernel}) / 10^9] / \text{time}$

- ▶ In GB/s

- ▶ Eg. Copy of 2048×2048 matrix (floats)

$$((2048^2 \times 4 \times 2) / 10^9) / \text{time}$$

Float: 4 bytes Read & write



Profiler bandwidth

- ▶ Not same as effective bandwidth, because:
 - ▶ Subset of GPU's multiprocessors used, and results extrapolated
 - ▶ Includes transfer of data not used by kernel
- ▶ Comparison of effective and profiler bandwidth can indicate how much bandwidth wasted by suboptimal coalescing