

# Parallel & Distributed Processing II:

## *parallel processing on manycore chips*

Nvidia GPUs and CUDA: Achieving High Performance

Eric Aubanel  
Winter 2010, UNB Fredericton

# Ideal CUDA programs

---

- ▶ High intrinsic parallelism
  - ▶ e.g. per-element operations
- ▶ Minimal communication (if any) between threads
  - ▶ limited synchronization
- ▶ High ratio of arithmetic to memory operations
- ▶ Few control flow statements
  - ▶ SIMD execution
    - ▶ divergent paths among threads in a block may be serialized (costly)

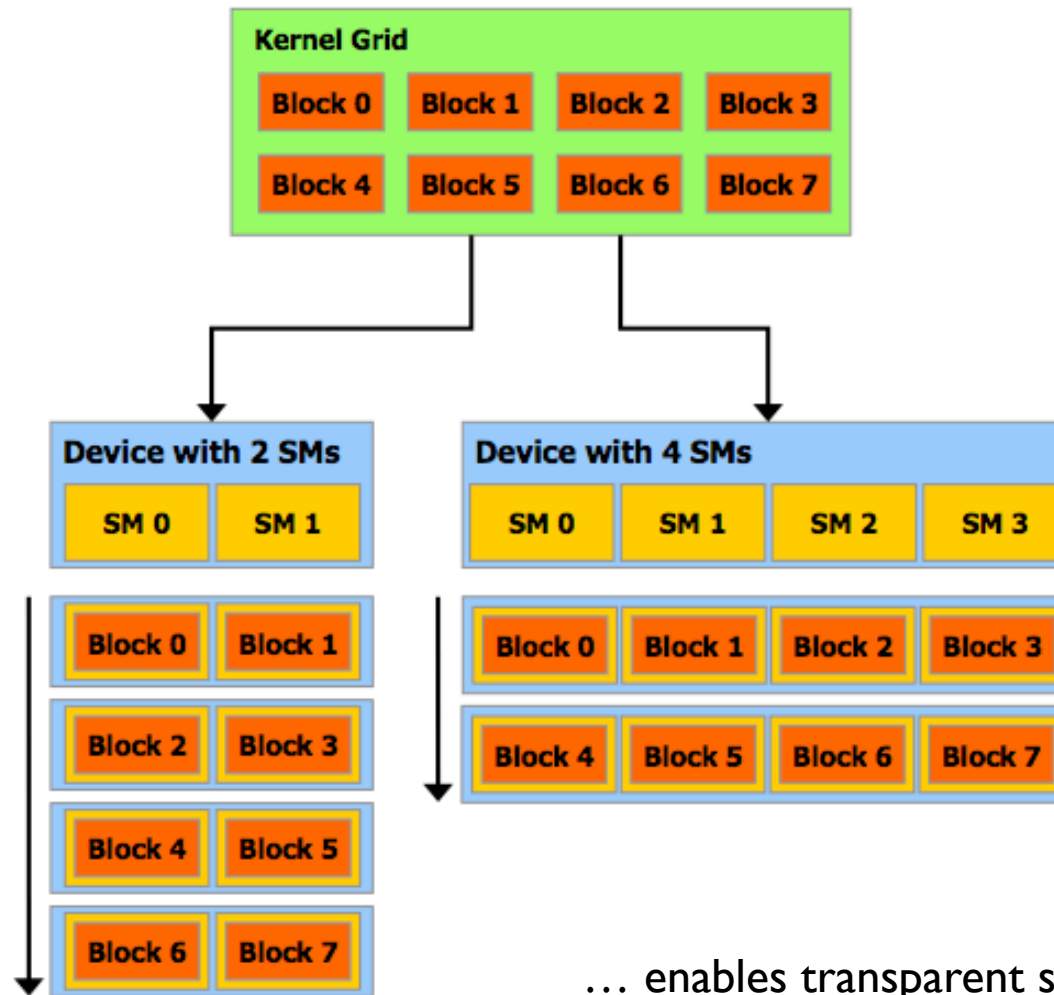
Source: John Mellor Crummey, Rice University

# Some Factors Affecting Performance

---


- ▶ Ratio of arithmetic to memory operations (arithmetic intensity)
- ▶ Conditional execution
- ▶ Maximum number of threads
  - ▶ Large enough to fully occupy all SMs
- ▶ Block size
  - ▶ Occupancy
- ▶ Efficient use of global memory
  - ▶ Enabling coalescing of memory accesses
- ▶ Threads sharing data
  - ▶ Use fast shared memory
- ▶ Efficient use of shared memory
  - ▶ Bank conflicts
- ▶ Host-device data transfer
- ▶ ...

# Block Scheduling



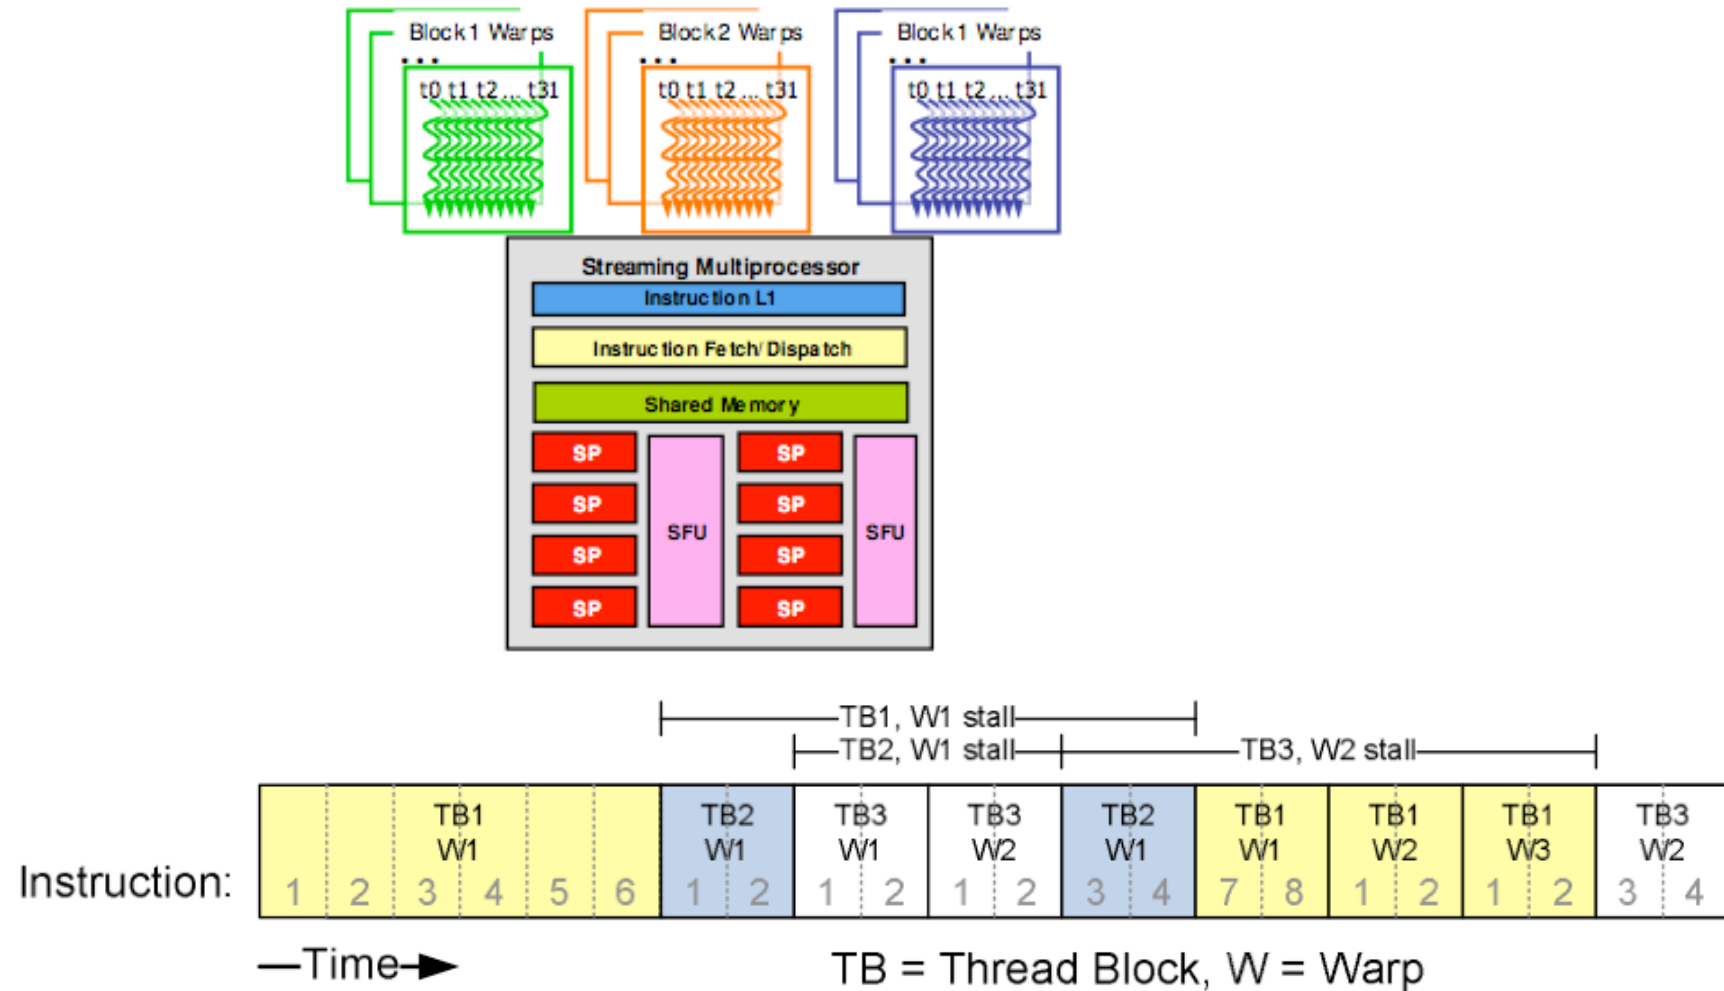
# Block Scheduling

- ▶ Maximum number of active blocks per SM: 8
  - ▶ As long as sufficient resources
    - ▶ Max of 768 threads
    - ▶ Registers, shared memory



Source: Wen-Mei Hwu, U of Illinois and David Kirk, Nvidia

# Blocks Divided into Warps



Source: Wen-Mei Hwu, U of Illinois and David Kirk, Nvidia

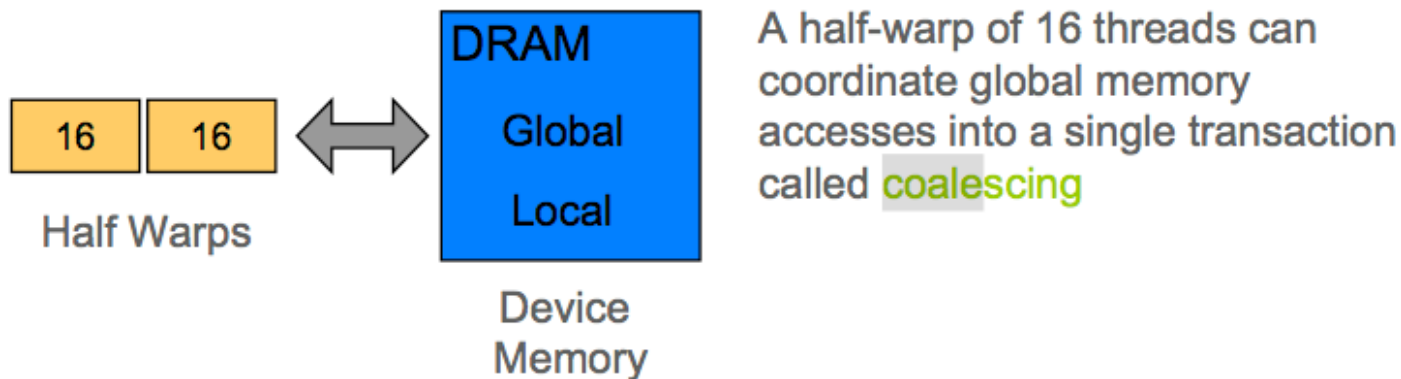
# Occupancy

---

- ▶  $(\# \text{ active warps/SM}) / (\max \# \text{ active warps} = 24)$ 
  - ▶ Blocks of 512 threads: only 1 block can be active at a time (since max 768 concurrent threads), so occupancy =  $16/24 = 2/3$ .
  - ▶ Blocks of 128 threads: 6 blocks can be active giving 768 threads and full (24/24) occupancy
- ▶ Full occupancy also limited by:
  - ▶ Registers
    - ▶ 32KB or 64KB per multiprocessor, partitioned among concurrent threads
  - ▶ Shared memory
    - ▶ 16KB per multiprocessor, partitioned among concurrent thread blocks
- ▶ See occupancy calculator
  - ▶ [http://developer.download.nvidia.com/compute/cuda/CUDA\\_Occupancy\\_calculator.xls](http://developer.download.nvidia.com/compute/cuda/CUDA_Occupancy_calculator.xls)

# Coalescing Global Memory Accesses

---




Dominik Göddecke, TU Dortmund



# Coalescing

- ▶ Global memory bandwidth used most efficiently when the simultaneous memory accesses by threads in a half-warp can be coalesced into a single memory transaction of 32, 64, or 128 bytes. Certain access requirements must be met
  - ▶ Depends on compute capability (see programming guide)
  - ▶ 1.0 and 1.1 have stricter access requirements


Examples – float (32-bit) data



Dominik Goddeke, TU Dortmund


# Coalescing

- ▶ Compute capability 1.0 and 1.1
  - ▶ K-th thread must access k-th word in the segment (or k-th word in two contiguous 128B segments for 128-bit words)
- ▶ Not all threads need to participate




Coalesces – 1 transaction

Out of sequence – 16 transactions



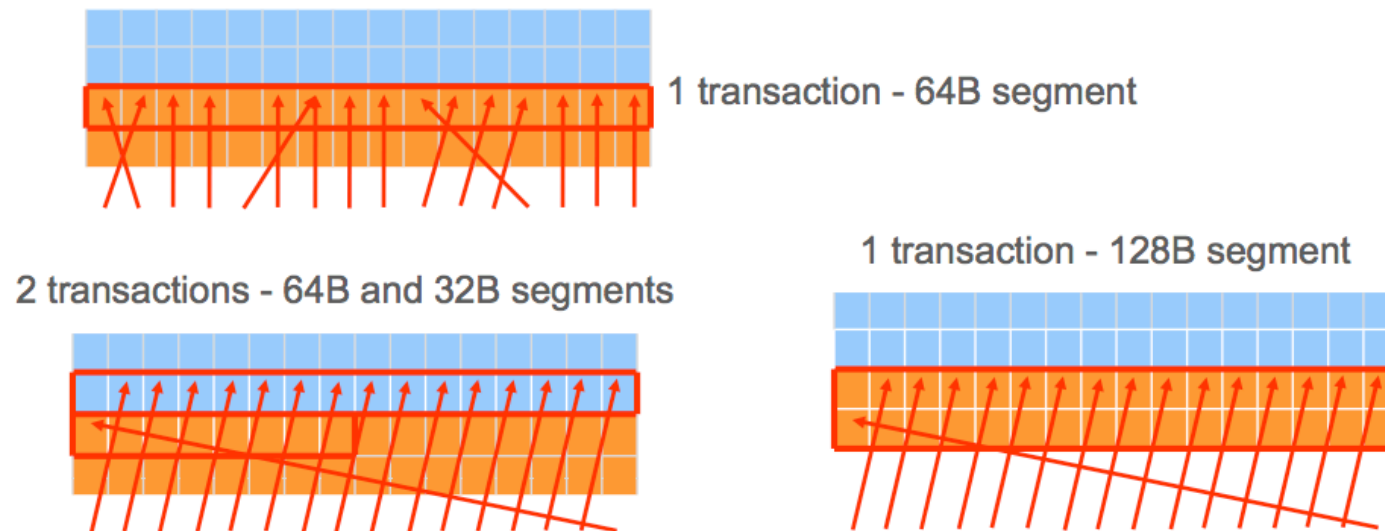
Misaligned – 16 transactions



Dominik Goddeke, TU Dortmund


# Coalescing

- ▶ Compute capability 1.2 and higher
  - ▶ Coalescing is achieved for any pattern of addresses that fits into a segment of size: 32B for 8-bit words, 64B for 16-bit words, 128B for 32- and 64-bit words
  - ▶ Smaller transactions may be issued to avoid wasted bandwidth due to unused words



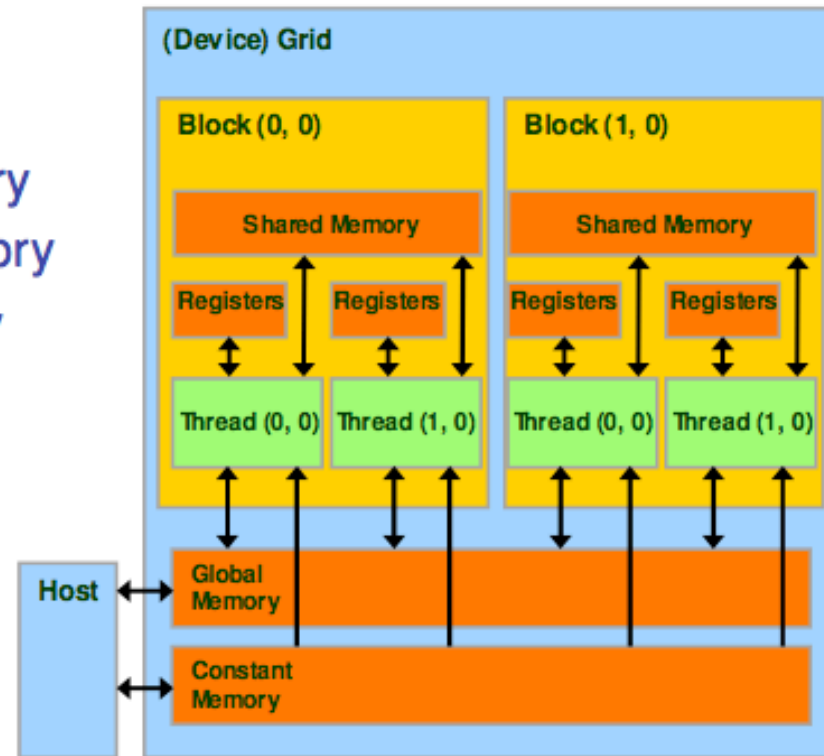
# Importance of Coalescing

```
__global__  
void saxpy_with_stride(int n, float a, float * x, float * y, int stride)  
{  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
  
    if (i < n)  
        y[i * stride] = a * x[i * stride] + y[i * stride];  
}
```



# Recall: CUDA Memory Model

- Device code can:
  - R/W per-thread **registers**
  - R/W per-thread **local memory**
  - R/W per-block **shared memory**
  - R/W per-grid **global memory**
  - Read only per-grid **constant memory**
- Host code can
  - R/W per grid **global** and **constant memories**



Source: Wen-Mei Hwu, U of Illinois and David Kirk, Nvidia

# Accessing CUDA memories

Variable declaration	Memory	Scope	Lifetime
Automatic variables other than arrays	register	thread	kernel
Automatic array variables	global	thread	kernel
<code>__device__ __shared__ int SharedVar;</code>	shared	block	kernel
<code>__device__ int GlobalVar;</code>	global	grid	application
<code>__device__ __constant__ int ConstVar;</code>	constant	grid	application

optional

Source: Wen-Mei Hwu, U of Illinois and David Kirk, Nvidia

# Shared Memory

---

- ▶ Each SM has 16 KB of Shared Memory
  - ▶ 16 banks of 32bit words
- ▶ CUDA uses Shared Memory as shared storage
  - ▶ visible to all threads in a thread block
- ▶ Useful when more than one thread in a block read same memory location
  - ▶ Reduce reads from global memory
- ▶ Can be used to take advantage of coalescing of global memory accesses
  - ▶ Coalesced copy to shared memory, followed by any pattern of access within shared memory
    - ▶ Except watch for bank conflicts

# Shared Memory Allocation

---

- ▶ 2 modes

- ▶ Static size within kernel

```
__shared__ float vec[256];
```

- ▶ Dynamic size when calling the kernel

```
// in main
```

```
int VecSize= 256 * sizeof(float);
```

```
foo<<<nBlocks, blockSize, VecSize>>>(p1, p2, ...);
```


```
// declare as extern within kernel
```

```
extern __shared__ float vec[];
```

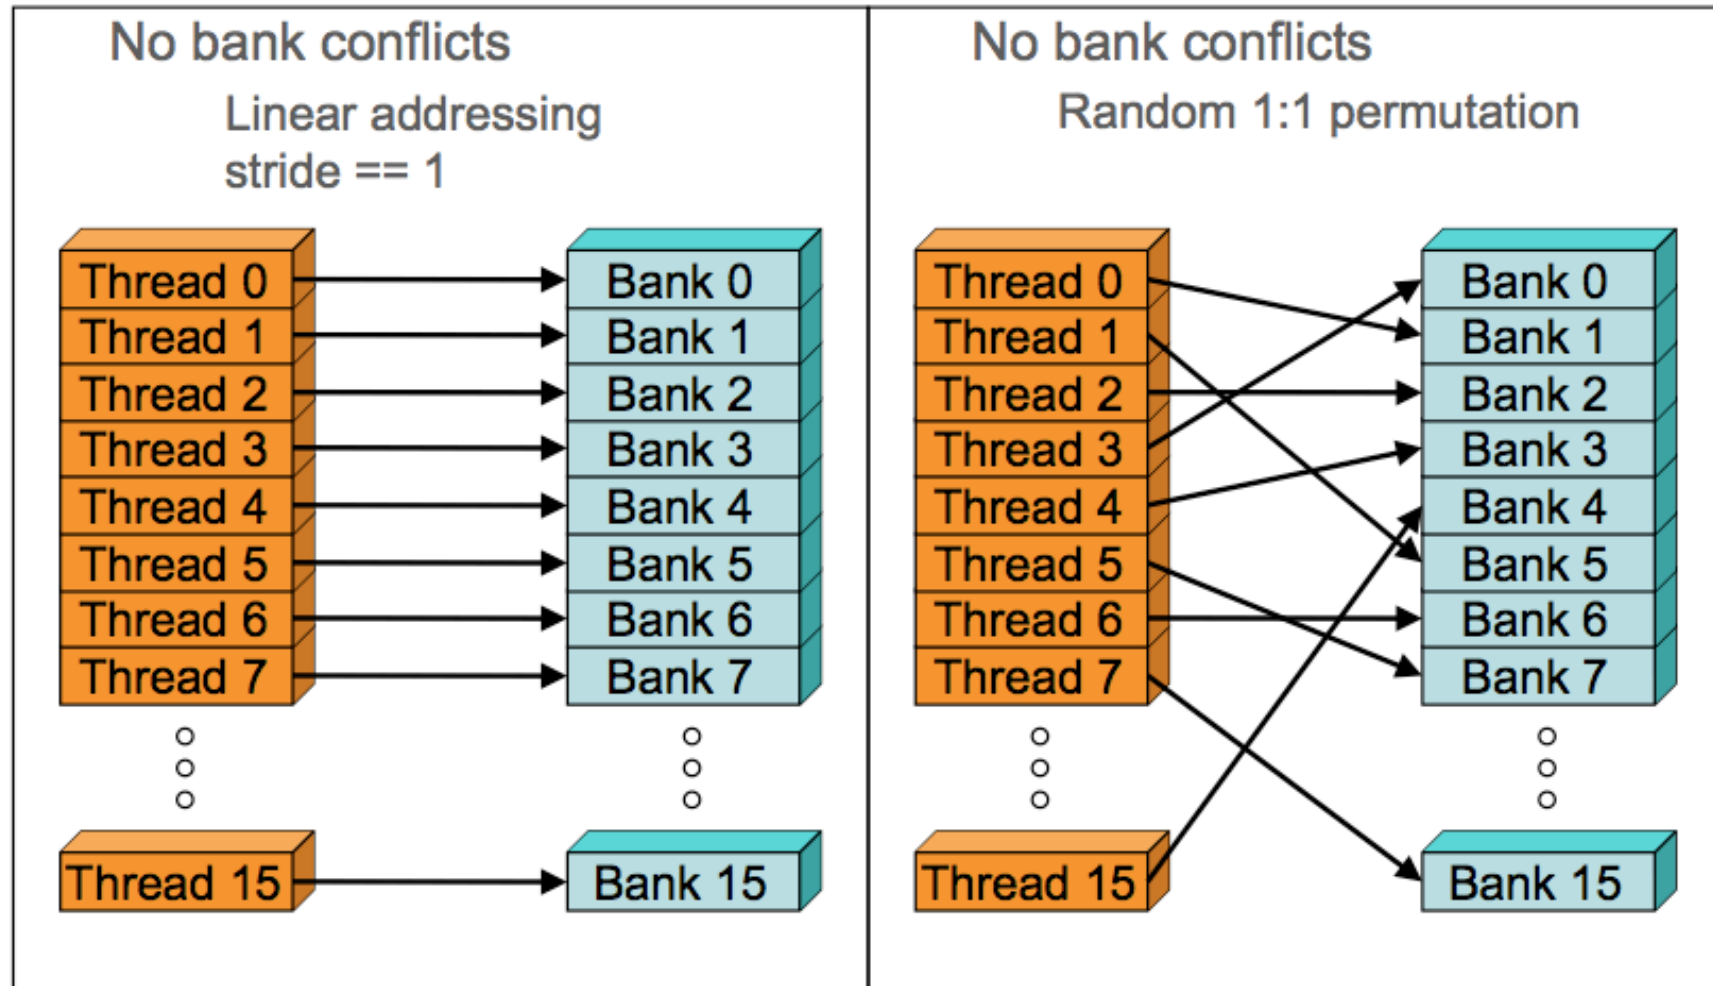


# Shared Memory Architecture

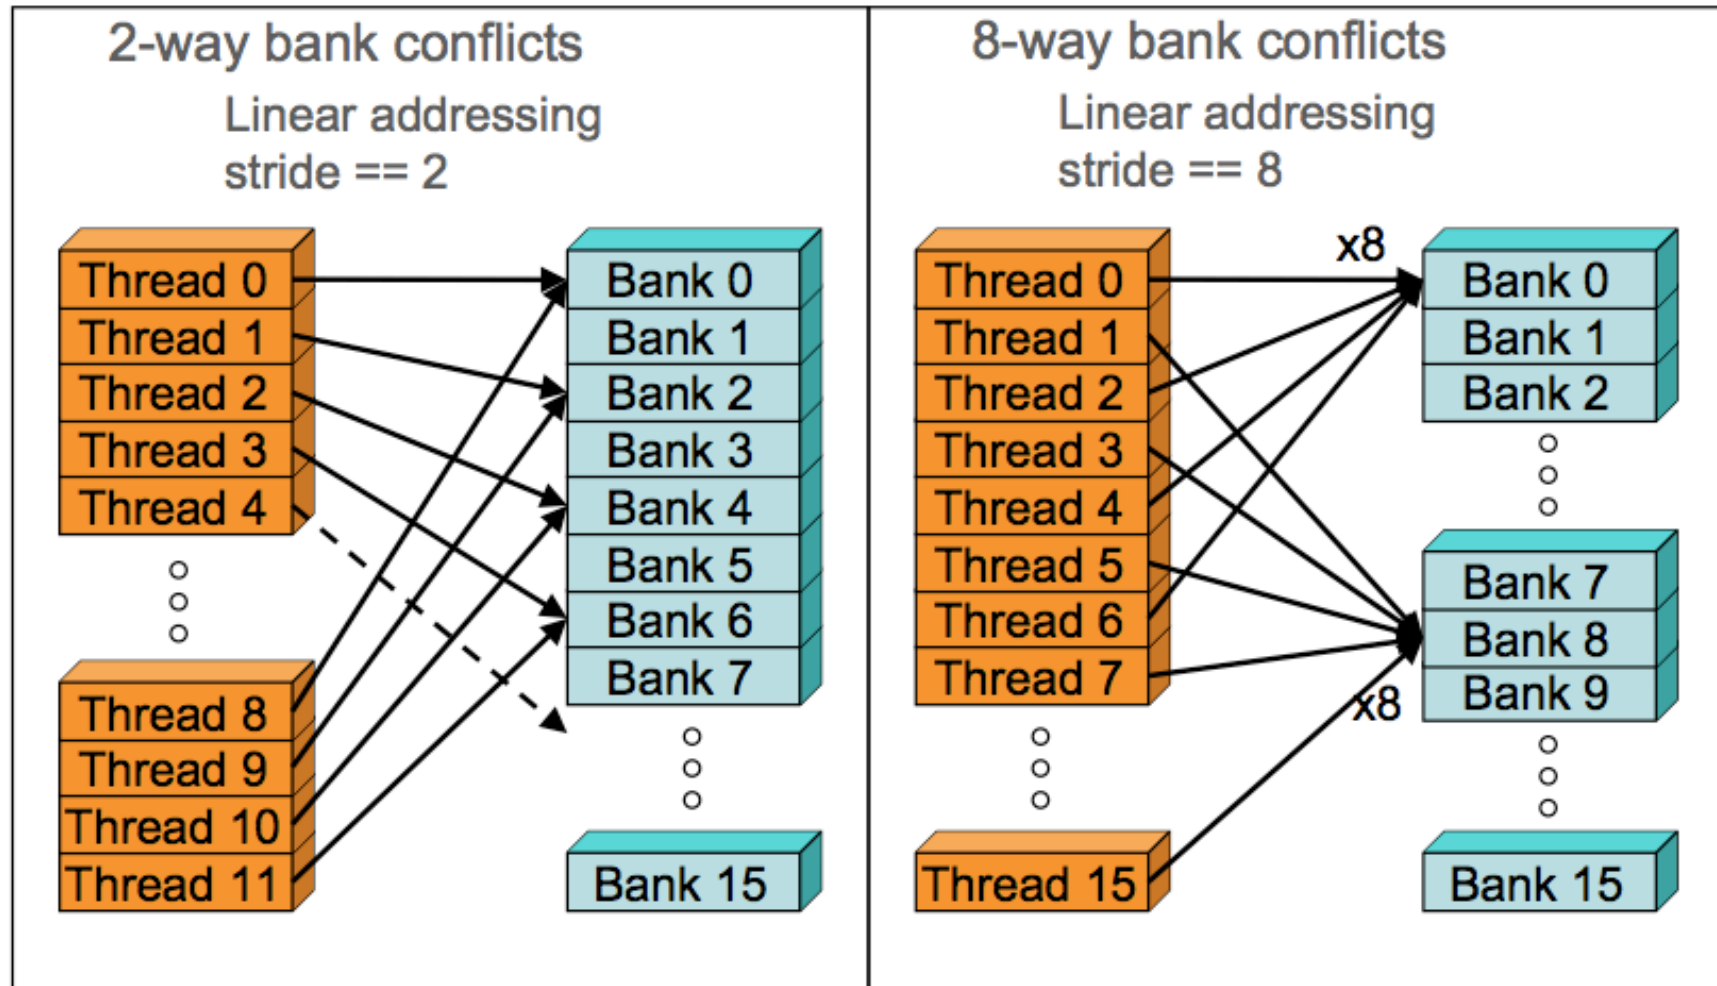
- ▶ Many threads accessing memory
  - ▶ Therefore, memory is divided in banks
  - ▶ Successive 32-bit words assigned to successive banks
- ▶ Each bank can serve one address per cycle
  - ▶ A memory can service as many simultaneous accesses as it has banks
- ▶ Multiple simultaneous accesses to a bank result in a bank conflict
- ▶ Conflicting addresses are serialized



# Bank Addressing Examples



# Bank Addressing Examples



# Bank Conflicts

---

- ▶ Shared memory is  $\sim$  as fast as registers
  - ▶ If there are no bank conflicts
  - ▶ `warp_serialize` profiler signal (see CUDA profiler)
- ▶ **The fast case**
  - ▶ If all threads of a half-warp access different banks, there are no bank conflicts
  - ▶ If all threads of a half-warp read the identical address, there is no bank conflict (broadcast)
- ▶ **The slow case**
  - ▶ Bank conflict: multiple threads in the same half-warp access the same bank
  - ▶ Must serialize the accesses
  - ▶ Cost = max # of simultaneous accesses to a single bank

# CUDA Profiler

---

- ▶ Accessible via GUI (cudaprof) and text interface
- ▶ Four environment variables for textual profiler:

**CUDA\_PROFILE:** Set to 1 or 0 to enable/disable the profiler

**CUDA\_PROFILE\_LOG:** Set to the name of the log file

(The default is ./cuda\_profile.log)

**CUDA\_PROFILE\_CSV:** Set to 1 or 0 to enable or disable a comma separated version of the log

**CUDA\_PROFILE\_CONFIG:** Specify a configuration file with up to four signals

# CUDA Profiler Signals

---

- ▶ **gld\_incoherent**: Number of non-coalesced global memory loads
- ▶ **gld\_coherent**: Number of coalesced global memory loads
- ▶ **gst\_incoherent**: Number of non-coalesced global memory stores
- ▶ **gst\_coherent**: Number of coalesced global memory stores
- ▶ **local\_load**: Number of local memory loads
- ▶ **local\_store**: Number of local memory stores
- ▶ **branch**: Number of branch events taken by threads
- ▶ **divergent\_branch**: Number of divergent branches within a warp
- ▶ **instructions**: instruction count
- ▶ **warp\_serialize**: Number of threads in a warp that serialize based on address conflicts to shared or constant memory
- ▶ **cta\_launched**: executed thread blocks

# Using the Profiler

---

- ▶ Profiler can only target one of the multiprocessors in the GPU
  - ▶ counter values will not correspond to the total number of warps launched for a particular kernel.
  - ▶ Best used to identify relative performance differences
- ▶ First set environment variables

```
export CUDA_PROFILE=1
export CUDA_PROFILE_CONFIG=
    $HOME/.cuda_profile_config
```
- ▶ Put up to four signals in `.cuda_profile_config`

```
gld_coherent
gld_incoherent
gst_coherent
gst_incoherent
```
- ▶ After running program, see profile data in log file.